




# **Linux Memory Management- kmalloc and vmalloc**

[www.embeddedshiksha.com](http://www.embeddedshiksha.com)

## **kmalloc:-**

The `kmalloc` is a kernel function used to obtain kernel memory in byte-sized chunks. If we need whole pages, then we have to use function mentioned in previous section.

**`void * kmalloc(size_t size, gfp_t flags)`**



## **Return Values:-**

On success: It returns memory at least size bytes in length.

On failure: Returns NULL.

So we have to check for return value after a call to kmalloc.

## **kfree**

`void kfree(const void *ptr)`

This method will free block of memory allocated by kmalloc.

## **gfp\_mask Flags**

We have seen that the low level page allocators and kmalloc expect gfp flags as a parameter. We will discuss about these flags further and conclude this article.

Here gfp stands for – get free pages and these gfp flags are formed using various other flags. Here our knowledge of zones will come in to picture.

The gfp flags are broken up into three categories: action modifiers, zonemodifiers, and types.

## **Action Modifiers:-**

Action modifiers specify how the kernel is supposed to allocate the requested memory. For example various contexts exist where developer needs to use different action modifiers:

- For Interrupt handlers must instruct the kernel not to sleep (because interrupt handlers cannot reschedule) in the course of allocating memory.
- In Process context – process can sleep in the course of allocating memory.

## **Low Level GFP Flags**

`__GFP_WAIT`

Indicates that the caller is not high priority and can sleep or reschedule.

`__GFP_HIGH`

The allocator can access emergency pools and Used by a high priority or kernel process.

**\_\_GFP\_IO**

Indicates that the caller can perform low level IO.

**\_\_GFP\_FS**

The allocator can start filesystem I/O.

**\_\_GFP\_HIGHIO**

Determines that IO can be performed on pages mapped in high memory

**Zone Modifiers** :-Zone modifiers specify from which memory zone the allocation should originate. Normally,allocations can be fulfilled from any zone.

**Zone Modifiers**

Flag Description

**\_\_GFP\_DMA** :- Allocates only from ZONE\_DMA

**\_\_GFP\_DMA32** :- All ocates only from ZONE\_DMA32

**\_\_GFP\_HIGHMEM**:-Allocates from ZONE\_HIGHMEM or ZONE\_NORMAL

**Type Flags:-**The type flags specify the required action and zone modifiers to fulfill a particular type of transaction.

Low Level GFP Flag Combinations for High Level

Flag	Modifier Flags
GFP_AUTOMATIC	_GFP_HIGH
GFP_NOIO	_GFP_WAIT
GFP_NOFS	{ _GFP_WAIT  _GFP_IO}
GFP_KERNEL	{ _GFP_WAIT  _GFP_IO  _GFP_FS}
GFP_USER	{ _GFP_WAIT  _GFP_IO  _GFP_FS}
GFP_HIGHUSER	{ _GFP_WAIT  _GFP_IO  _GFP_FS  _GFP_HIGH MEM}
GFP_DMA	_GFP_DMA

Kernel provides various high level flags used according to the context you are calling the `kmalloc` function:

## **GFP\_ATOMIC:-**

Used to allocate memory in interrupt handlers and other code outside of a process context and it never sleeps.

Interrupts handles are not executing in process context and they cannot call 'schedule' function. Also interrupt handles should execute for short period of time and so should not sleep.

## **GFP\_KERNEL**

Normal allocation of kernel memory. May sleep. This can be used for example, when executing a system call in kernel on behalf of a process. This can block.

## **GFP\_USER**

Used to allocate memory for user-space pages; it may sleep.

## **GFP\_HIGHUSER**

Like GFP\_USER, but allocates from high memory, if any.

GFP\_NOIO

GFP\_NOFS

A GFP\_NOFS allocation is not allowed to perform any filesystem calls, while GFP\_NOIO disallows the initiation of any I/O at all. They are used primarily in the filesystem and virtual memory code where an allocation may be allowed to sleep.




## Vmalloc

It allocates memory which is virtually contiguous and not necessarily physically contiguous. But do remember 'kmalloc' will return physically contiguous and virtual contiguous memory.

So if memory is virtually contiguous and physically not, then how virtual memory is mapped to physical memory?

Kernel does this by using page tables. This we discuss later articles.



You may also get a doubt, which is faster in allocating memory?  
The answer lies in the above discussion.

Since vmalloc use page tables to map the virtual address to physical address, so whenever you access memory allocated using vmalloc, you have to traverse through page table. This is additional overhead compared to kmalloc. So vmalloc should be used where it is really necessary.

`void * vmalloc(unsigned long size)`

**On success:-** Returns a pointer to at least size bytes of virtually contiguous memory.

**On failure:-** the function returns NULL.

To free an allocation obtained via vmalloc(), use

`void vfree(const void *addr)`



**Thank You**