

Memory Management 101: Introduction to Memory Management in Linux

Christopher Lameter, Ph.D. <cl@linux.com>



Jump Trading LLC

Overview

- ❑ Memory and processes
- ❑ Real/Virtual memory and Paging
- ❑ Machine configuration
- ❑ Processes use of memory
- ❑ Overcommit
- ❑ Knobs
- ❑ There is an advanced MM talk tomorrow called *“Flavors of Memory”*



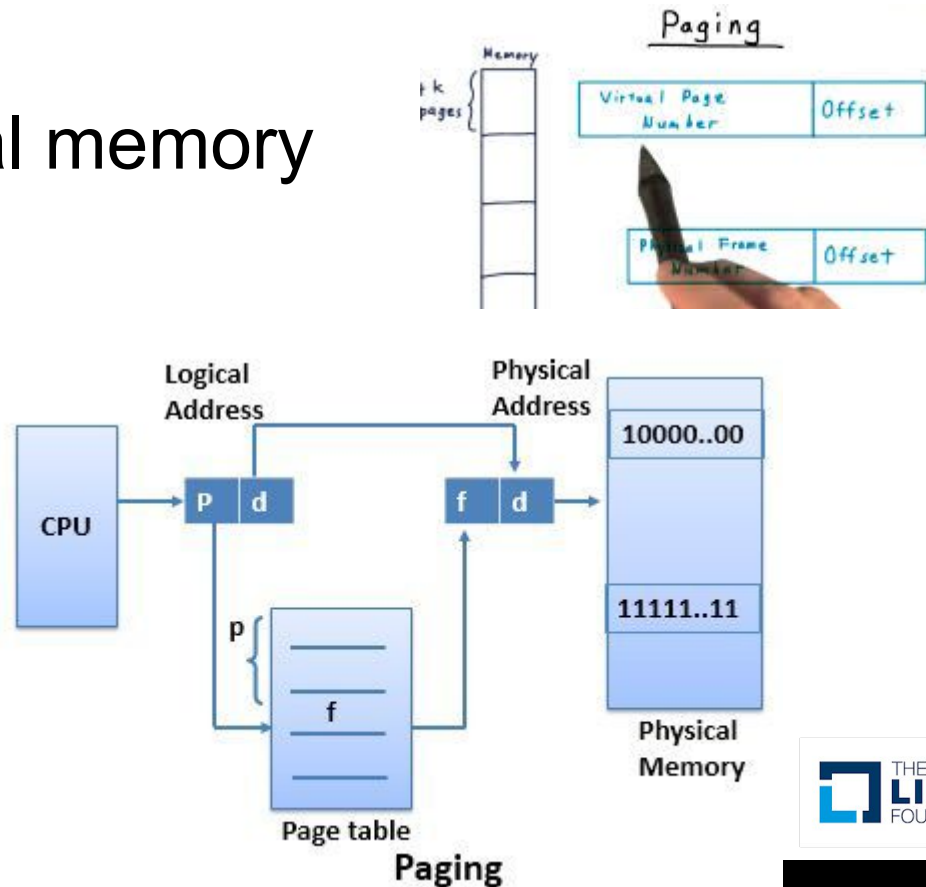
Pages and physical page frame numbers

- Division of memory into “pages”
 - 1-N bytes become split at page size boundaries and become
 $M = N / \text{page size}$ pages
- Refer to memory by the Page Frame Number (PFN) and an offset into the page.
- Common size is 4k (Intel legacy issues)
- MMU creates virtual addresses.




Basics of “paging”

- Process have virtual memory
- -> PFN
- Page Tables
- Faults
 - Major
 - Minor
- Virtual vs physical



Process memory


- Virtual memory maps to physical memory
- View of memory that is distinct for each process.
- Pages shared
- Access control
- Copy on Write



Swap, Zero pages etc.

- ❖ Swap page
- ❖ Zero page
- ❖ Read data behavior
- ❖ Write data behavior
- ❖ Anonymous vs file backed pages

I see a number zero!



Linux Basic memory information

`/proc/meminfo`

`/sys/devices/system/`
has lots of more detailed
information on hardware
(processors and
memory)

Commands:

`numactl --hardware`

`free, top, dmesg`

MemTotal:	31798552 kB
MemFree:	25949124 kB
MemAvailable:	30823580 kB
Buffers:	220988 kB
Cached:	4679188 kB
SwapCached:	0 kB
Active:	2803000 kB
Inactive:	2336992 kB
Active(anon):	240776 kB
Inactive(anon):	6432 kB
Active(file):	2562224 kB
Inactive(file):	2330560 kB
Unevictable:	0 kB
Mlocked:	0 kB
SwapTotal:	2097148 kB
SwapFree:	2097148 kB
Dirty:	48 kB
Writeback:	0 kB



AnonPages:	239716 kB
Mapped:	195596 kB
Shmem:	7396 kB
Slab:	550628 kB
SReclaimable:	443040 kB
SUnreclaim:	107588 kB
KernelStack:	6840 kB
PageTables:	11176 kB

Inspecting a process memory use

`/proc/<pid>/status`

`/proc/<pid>/*maps`

(there are other files in `/proc/<pid>/` with more information about the processes)

Commands:

ps, top



Name:	sshd
VmPeak:	65772 kB
VmSize:	65772 kB
VmLck:	0 kB
VmPin:	0 kB
VmHWM:	6008 kB
VmRSS:	6008 kB
RssAnon:	1216 kB
RssFile:	4792 kB
RssShmem:	0 kB

VmData:	1332 kB
VmStk:	132 kB
VmExe:	492 kB
VmLib:	8076 kB
VmPTE:	168 kB
VmSwap:	0 kB

User limit (ulimit)

- max memory size
- virtual memory
- stack size
- and lots of other controls.

```
cl@nuc-kabylake:/proc/6713$ ulimit -a
core file size          (blocks, -c)  0
data seg size           (kbytes, -d) unlimited
scheduling priority      (-e)  0
file size                (blocks, -f)  unlimited
pending signals          (-i)  123132
max locked memory        (kbytes, -l) 16384
max memory size          (kbytes, -m) unlimited
open files               (-n)   1024
pipe size                (512 bytes, -p) 8
POSIX message queues     (bytes, -q) 819200
real-time priority       (-r)  0
stack size               (kbytes, -s)  8192
cpu time                 (seconds, -t) unlimited
max user processes       (-u)  123132
virtual memory           (kbytes, -v) unlimited
file locks               (-x)  unlimited
```



Overcommit configuration

Virtual memory use vs physical

`overcommit_kbytes`

`overcommit_memory`

0 - overcommit. Guess if mem is available.

1 - Overcommit. Never say there is no memory

2 - Only allocate according to the ratio

`overcommit_ratio`

$\text{total} = \text{swap} + \text{physical} * \text{ratio}$



Important VM control knobs

Found in `/proc/sys/vm`

More descriptions of these knobs in Kernel source code. **linux/Documentation/admin-guide**

`admin_reserve_kbytes` **`dirty_writeback_centisecs`** `min_free_kbytes` `numa_zonelist_order`
`stat_refresh` `block_dump` `drop_caches` `min_slab_ratio` `oom_dump_tasks` `swappiness`
`compact_memory` `extfrag_threshold` `min_unmapped_ratio` `oom_kill_allocating_task`
`user_reserve_kbytes` `compact_unevictable_allowed` `hugetlb_shm_group` `mmap_min_addr`
`overcommit_kbytes` `vfs_cache_pressure` **`dirty_background_bytes`** `laptop_mode`
`mmap_rnd_bits` `overcommit_memory` `watermark_scale_factor` **`dirty_background_ratio`**
`legacy_va_layout` `mmap_rnd_compat_bits` `overcommit_ratio` `zone_reclaim_mode` `dirty_bytes`
`lowmem_reserve_ratio` `nr_hugepages` `page-cluster` `dirty_expire_centisecs` `max_map_count`
`nr_hugepages_mempolicy` `panic_on_oom` **`dirty_ratio`** `memory_failure_early_kill`
`nr_overcommit_hugepages` `percpu_pagelist_fraction` **`dirtytime_expire_seconds`**
`memory_failure_recovery` `numa_stat` `stat_interval`

Resources

- Admin Guide online
<https://www.kernel.org/doc/html/v4.14/admin-guide/index.html>
- Kernel.org has wikis and documentation
(www.kernel.org)
- manpages (especially for system calls and coding)



Questions / Comments

You can reach me at cl@linux.com or
[@qant](#) on twitter



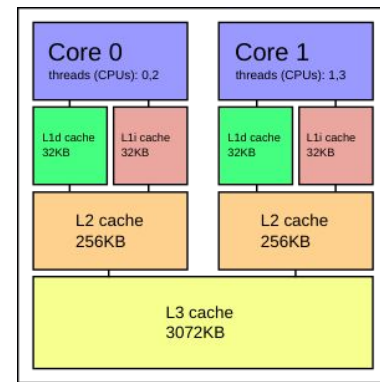
OPEN SOURCE SUMMIT

EUROPE


THE LINUX FOUNDATION

“Simple” Memory Access

- **UMA** (Uniform Memory Access)
- Any access to memory has the same characteristics (performance and latency)
- The vast major of systems have only UMA.
- But there is always a processor cache hierarchy
 - The CPU is fast, memory is slow
 - Caches exist to avoid accesses to main memory
- **Aliasing**
- **Coloring**
- Cache Miss
- Trashing





CPU Cache Access Latencies in Clock Cycles



NUMA Memory

- Memory with different access characteristics
- Memory **Affinities** depending on where a process was started
- Control **NUMA** allocs with memory policies
- System Partitioning using Cpusets and Containers
- Manual memory **migration**
- Automatic memory migration




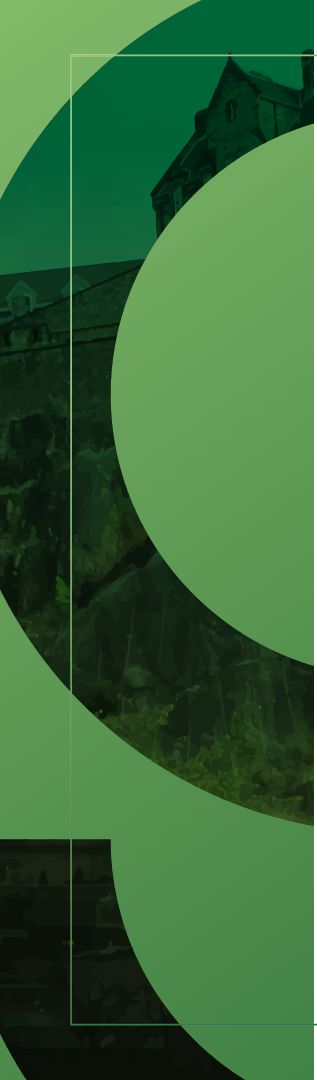
Huge Memory

- Typical memory is handled in chunks of base page size (Intel 4k, IBM PowerX 64K, ARM 64K)
- Systems support larger memory chunks of memory called Huge pages (Intel 2M)
- Must be pre configured on boot in order to guarantee that they are available
- Required often for I/O bottlenecks on Intel.
- 4TB requires 1 billion descriptors with 4K pages.

Most of this is needed to compensate for architectural problems on Intel. Intel processors have difficulties using modern SSDs and high speed devices without this.

- Large contiguous segments (I/O performance)
- Fragmentation issues
- Uses files on a special file system that must be explicitly requested by mmap operations from special files.





An Introduction to Linux memory management. The basics of paging. Understanding basic hardware memory management and the difference between virtual, physical and swap memory. How do determine hardware installed and how to figure out how processes use that memory. How a process uses physical and virtual memory effectively. How to control overcommit and virtual and/or physical memory limits.

Basic knobs in Linux to control memory management. System calls for a process to control its memory usage