

CS162

Operating Systems and Systems Programming

Lecture 14

Virtual Memory


Professor Natacha Crooks

<https://cs162.org/>

Slides based on prior slide decks from David Culler, Ion Stoica, John Kubiatawicz,
Alison Norman and Lorenzo Alvisi

Please fill in Midterm Survey
Constructive feedback please!
Help us make the class better

Topic Breakdown



Recall: A process

A process is an **instance** of a running program

CPU

Memory
(address space)

Registers

IO
information

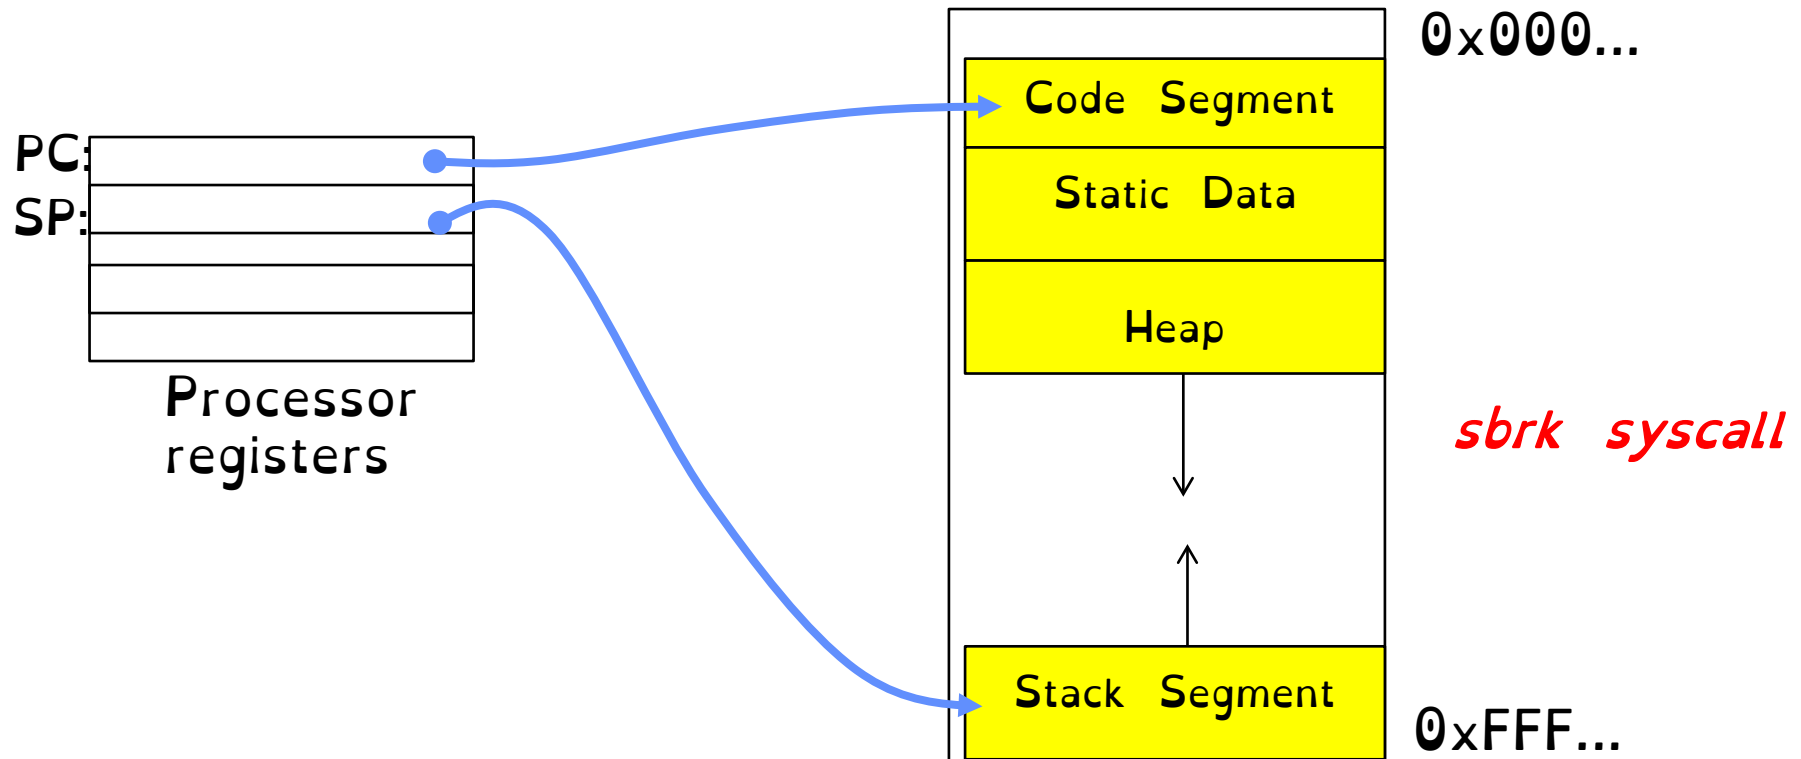
Store code,
data, stack,
heap

Program
Counter,
Stack
Pointer
Regular
registers

Open files
(and others)

Recall: Address Space

Set of memory addresses accessible to program
(for read or write)



Memory Virtualization Objectives

Isolation

Flexibility

Infinite Resources

How can we do so **efficiently**?

Interposing on Process Behaviours

OS interposes on process's IO operations
via Syscalls

OS interposes on process's CPU usage
Via Preemption

How can OS interpose on process's memory access?

Too slow for the OS to interpose every memory access.
Translation: hardware support to accelerate common case.
Uncommon cases “trap” into the OS to handle

An Address

A memory address refers to the location of a byte in memory.

Most machines are byte-addressable



K bits



2^K
things

Bits & Addresses

If an address space has 32 bits,
how many unique addresses do I have?

$$2^{32} = (4294967296)$$

$$2^{64} = \text{more than the atoms of the universe}$$

How many bits necessary to exclusively enumerate 4 elements?

$$2 \text{ bits} \Rightarrow 2^2 = 4. \Rightarrow \log_2(4)$$

How many 32 bit numbers fit in a 2^{32} address space?

$$32 \text{ bits} \rightarrow 4 \text{ bytes} \rightarrow 2^2.$$

$$2^{32}/2^2 = 2^{30}, 1 \text{ billion}$$

Increasingly powerful mechanisms

No protection. Living
life on the edge



Base & Bound

Base & Bound with
Relocation


Segmentation

Paging

Uniprogramming: I'm all alone

Application always runs at same place in physical memory
since only one application at a time

Application can access any physical address




Application given illusion of dedicated machine by giving it
reality of a dedicated machine

Memory Translation Through Relocation

Use loader/linker to adjust addresses when program loaded into memory.


Memory Translation Through Relocation



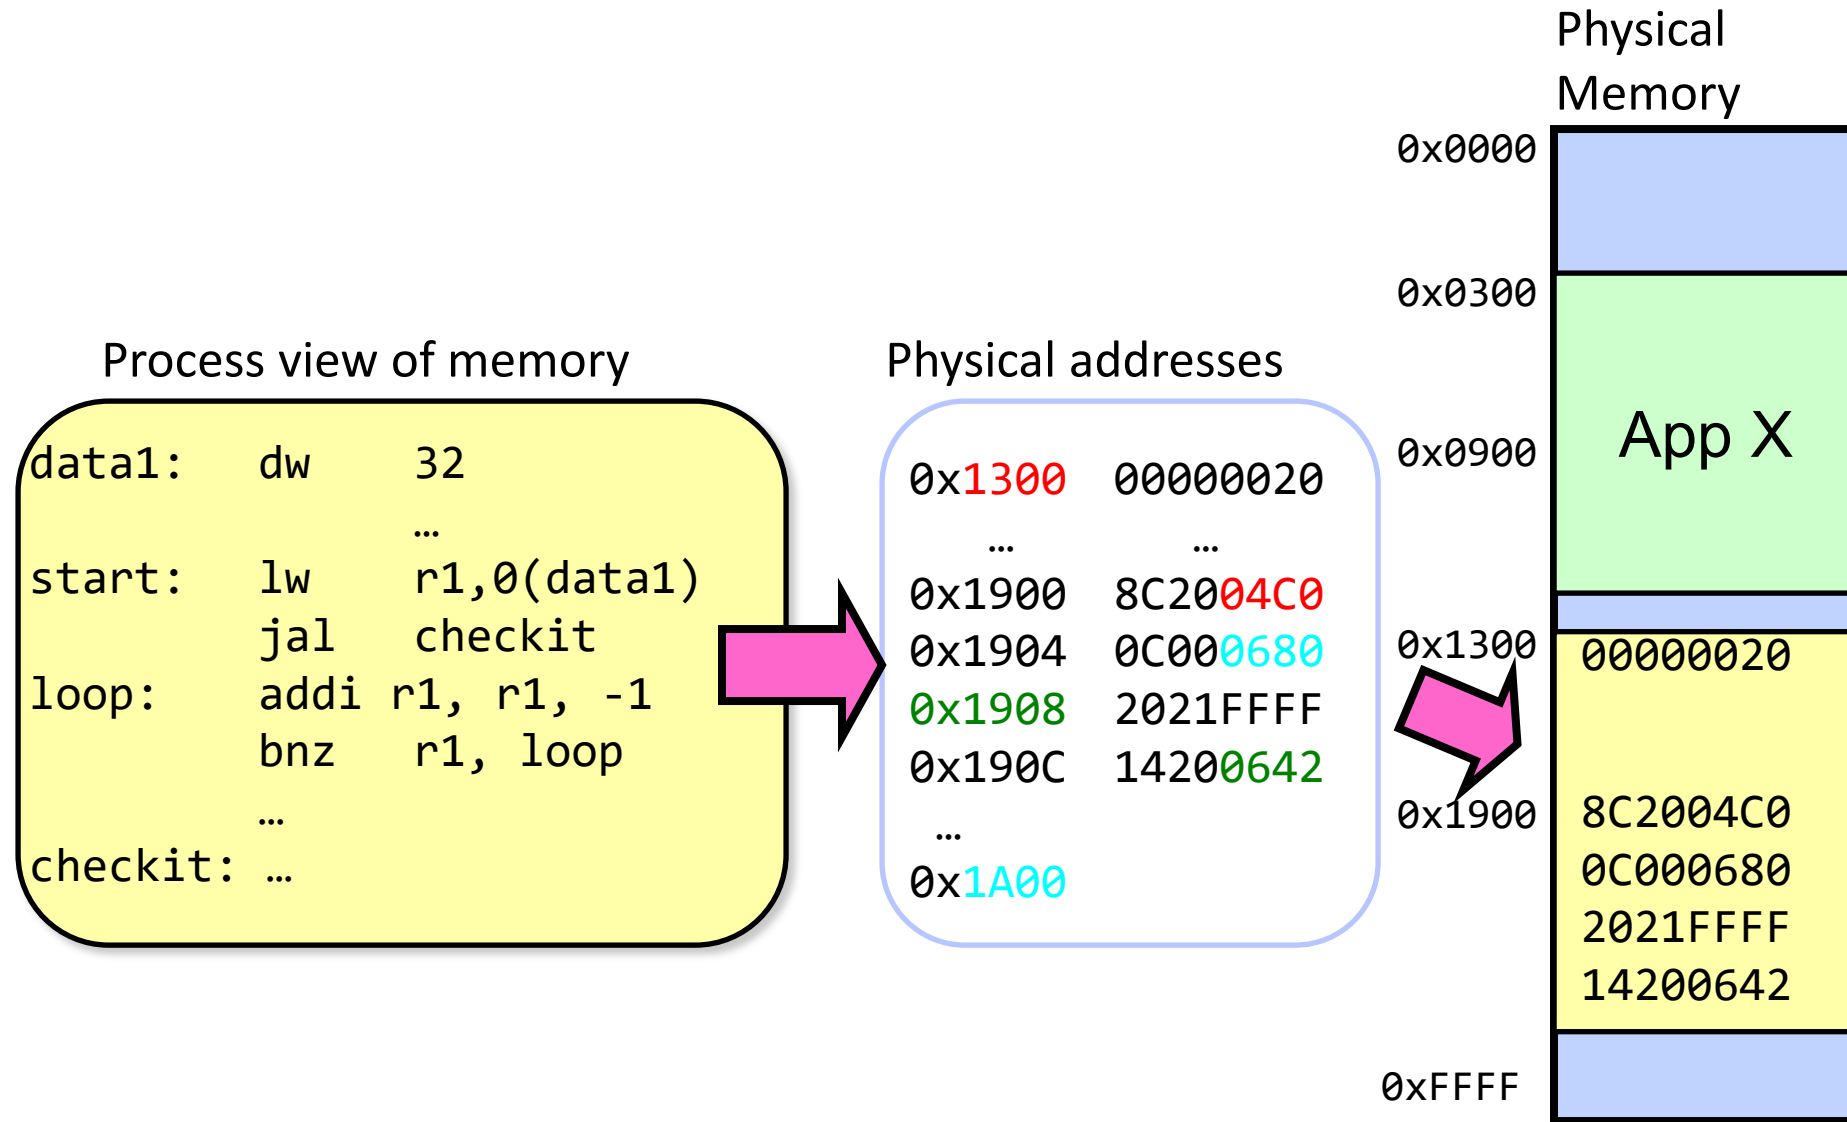
0xFFFFFFFF

0x00020000

0x00000000




Memory Translation Through Relocation



Memory Translation Through Relocation

With this solution, no protection: bugs in any program can cause other programs to crash or even the OS



Recall: A Bug's Tail

The character could leave the game area and start overwriting other running programs and kernel memory.

One of the worst bugs I ever had to deal with was in this game. Once the game player made it to the Colony, every so often the system would crash and burn at totally random times. You might be playing for ten minutes when it happened or ten hours, but it would just die in a totally random way

There was a slow-moving slug like creature that knew how to follow the game player's trail. When it came across another creature, rather than bouncing off and risk losing the trail, I made it so that it would destroy the other creature and stay on target to find you. This worked great, except that on some rare occasions, this slug could do to a wall what it did to the other creatures. That is, it could delete it. This meant that the virtual door was now open for this creature to explore the rest of the RAM on the Macintosh, deleting and modifying it as it went along. Of course, it was just a matter of time before it found some juicy code. In other words, the bug was a REAL bug.

Recall: Super Mario Land 2

Mario could exit a level and explore the entire memory of the system



Increasingly powerful mechanisms

No protection. Living
life on the edge



Base & Bound

Base & Bound with
Relocation

Segmentation

Paging


Recall: Memory Protection

OS and applications both resident in memory

Application should not read/write kernel memory
(or other apps memory)


Base And Bound

Hardware to the rescue!
Base and Bound registers



Base & Bound

Hardware to the rescue!
Base and Bound registers



Base & Bound

Kernel Mode executes without
Base and Bound registers

Loader rewrites address to the desired offset in physical
memory.

Relocation

`movl 1000, %eax`

`movl 4000, %eax`

Limitations of Base & Bound

1) No expandable memory

Static memory allocation

2) No memory Sharing

Cannot share memory between processes

3) Non-Relative Memory Addresses

Location of code & data determined at runtime

4) External Fragmentation

Cannot relocate/move programs. Leads to fragmentation

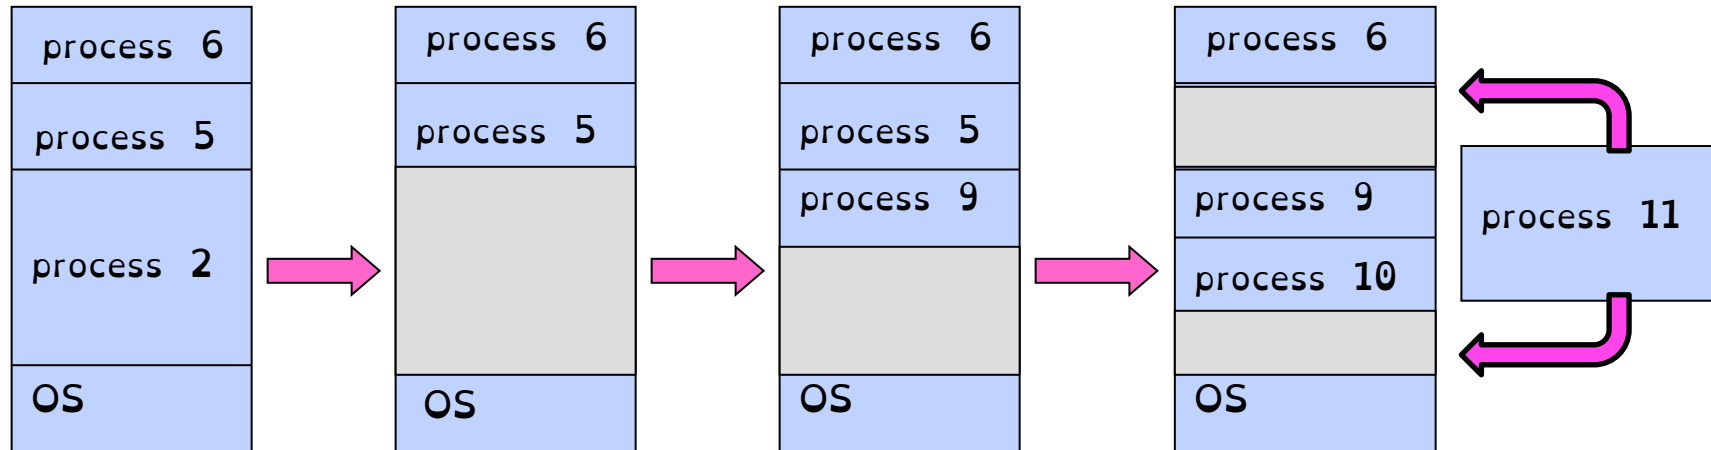
5) Internal Fragmentation

Address Space must be contiguous

Fragmentation in More Detail

External Fragmentation


Free chunks between allocated regions



Fragmentation in More Detail

Internal Fragmentation

Space inside allocated address space may not be fully used.



Limitations of Base & Bound

1) No expandable memory

Static memory allocation

2) No memory Sharing

Cannot share memory between processes

3) Non-Relative Memory Addresses

Location of code & data determined at runtime

4) External Fragmentation

Cannot relocate/move programs. Leads to fragmentation

5) Internal Fragmentation

Address Space must be contiguous

Increasingly powerful mechanisms

No protection. Living
life on the edge

Base & Bound

Base & Bound with
Relocation

Segmentation


Paging

Base & Bound With Hardware Relocation

Address Translation

Virtual address space

Set of memory addresses that
process can “touch”



Physical address space


Set of memory addresses
supported by hardware

Base And Bound With Relocation

Each program is written and compiled as if it is loaded at address zero


Memory references are **translated** by the processor

$\text{physical address} = \text{virtual address} + \text{base}$



Memory Management Unit

Hardware that performs translation of virtual to physical addresses



Limitations of Base & Bound with Relocation

1) No expandable
memory

Static memory
allocation

2) No memory Sharing

Cannot share memory
between processes

3) Non-Relative Memory
Addresses

Location of code &
data determined at
runtime

4) External
Fragmentation

Cannot relocate/move
programs. Leads to
fragmentation

5) Internal
Fragmentation

Address Space must be
contiguous

Increasingly powerful mechanisms

No protection. Living
life on the edge

Base & Bound

Base & Bound with
Relocation

Segmentation

Paging


Segmentation

Create a base and bounds pair per **logical segment** of the address space

A **segment** is a contiguous portion of the address space of a particular length

Can place each segment **independently** at different locations in memory

Segmentation



Minimises internal fragmentation
(code, data, heap, stack segments placed independently)

Implementation of a multi-segment model

Segment map resides in processor


Segment number mapped into base/limit pair

Base added to offset to generate physical address

0:	Base0	Limit0
1:	Base1	Limit1
2:	Base2	Limit2
3:	Base3	Limit3
4:	Base4	Limit4
5:	Base5	Limit5
6:	Base6	Limit6
7:	Base7	Limit7
8:	Base7	Limit7

Address Translation

A logical address consists of two parts: a segment identifier (top bits) and an offset that specifies the relative address within the segment (bottom bits)



Address Translation

Assume we have 16 bit addresses

Question: if I have 4 segments (code, data, stack, heap),
how many segment bits do I need?

$$\text{Log}(4) = 2$$

Segment	0:	00
Segment	1:	01
Segment	2:	10
Segment	3:	11

Address Translation

Assume we have 16 bit addresses

Question: if I have 4 segments (code, data, stack, heap),
how many segment bits do I need?

$$\text{Log}(4) = 2$$

Question: what is the maximum size of each segment?

$$16 - 2 = 14 \text{ bits left. } \Rightarrow 2^{14} \text{ bytes}$$

Question: if I have 7 segments and an address size of
32 bits, what is the maximum size of a segment?


$$\text{Log}_2(7) = 2.8 \Rightarrow 3 \text{ bits. } 2^{(32-3)} = 2^{29}$$

Example: Four Segments (16 bit addresses)




Virtual Address Format

Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000



Virtual
Address Space



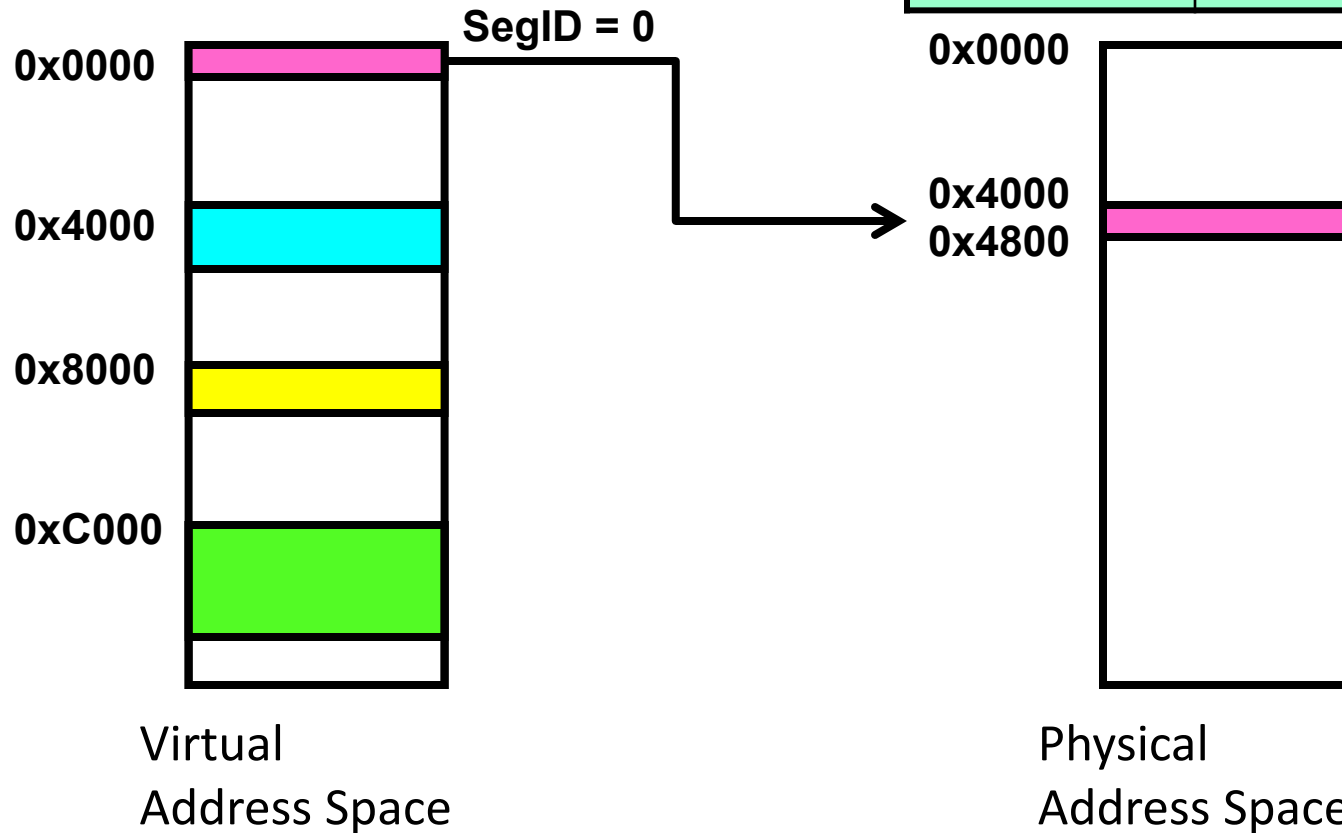
Physical
Address Space

Example: Four Segments (16 bit addresses)



Virtual Address Format

Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000

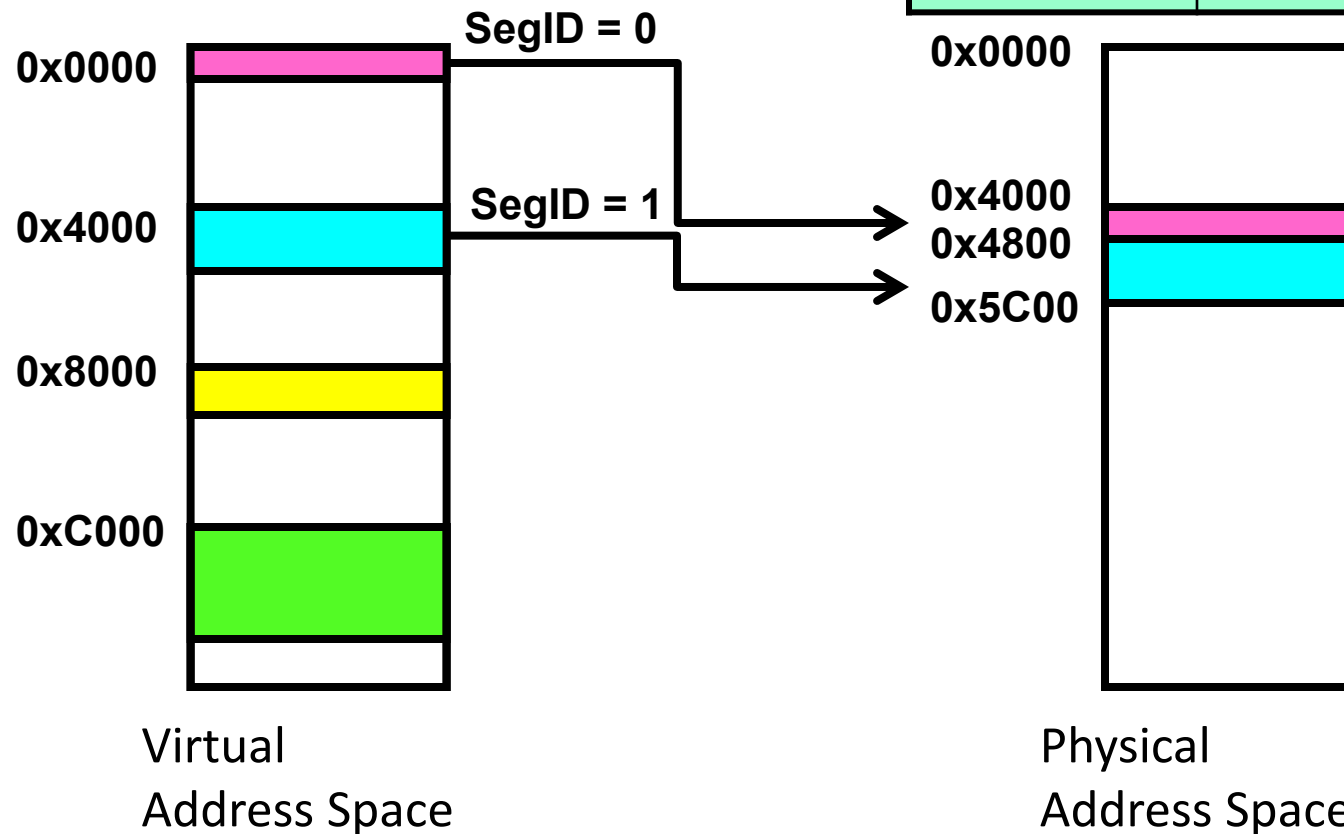


Example: Four Segments (16 bit addresses)



Virtual Address Format

Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000

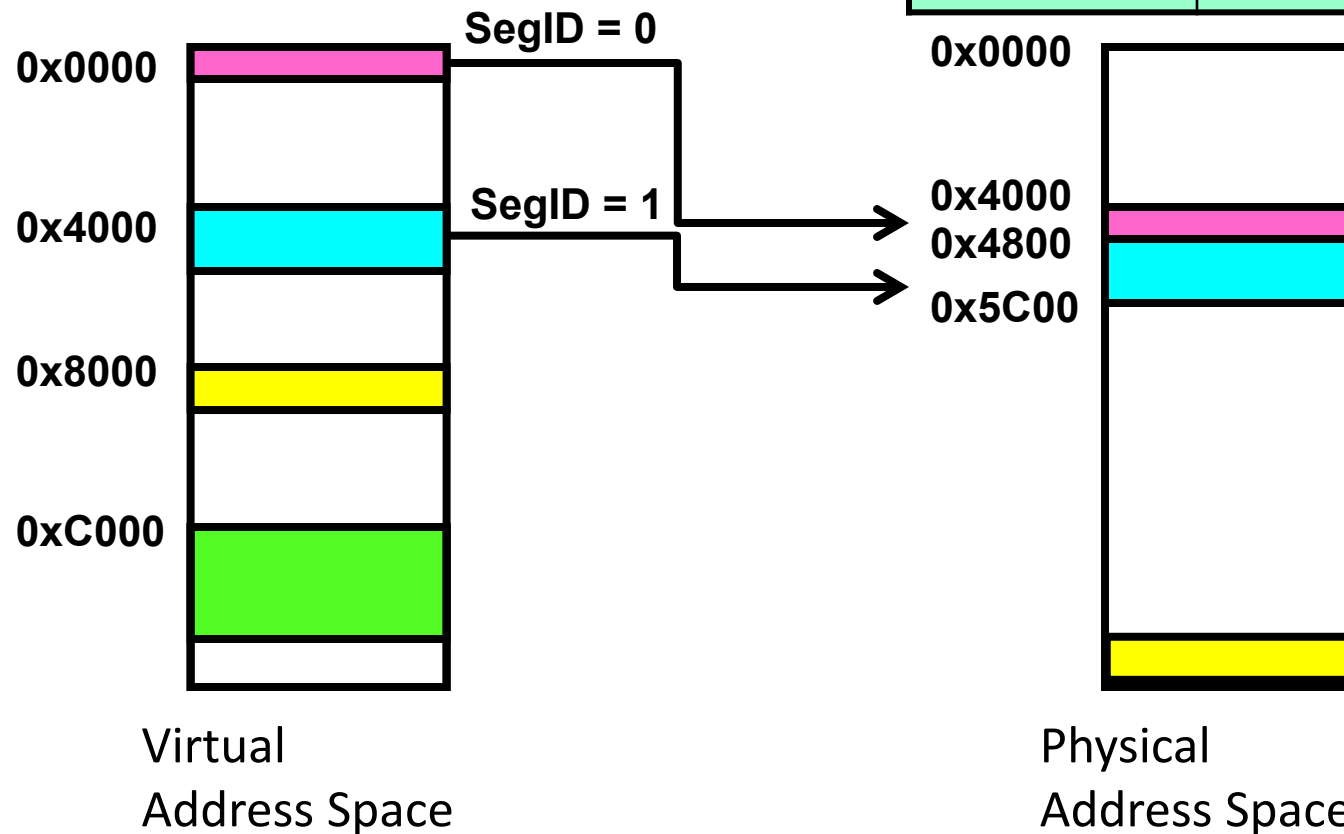


Example: Four Segments (16 bit addresses)



Virtual Address Format

Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000



Example: Four Segments (16 bit addresses)



Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000

0x4000

01 00 0000 0000 0000

Segment 1 Offset 0x0 0x4000

0x8020

10 00 0000 0010 0000

Segment 2 Offset 0x20 0xF020

Adding support for sharing

Useful to share
certain memory segments between address spaces.

Seg ID #	Base	Limit	Protection Bits
0 (code)	0x4000	0x0800	Read-Execute
1 (data)	0x4800	0x1400	Read-Write
2 (shared)	0xF000	0x1000	Read-Write
3 (stack)	0x0000	0x3000	Read-Write

Hardware must now check whether access is
1) within bounds 2) permissible

Segmentation Summary Pros

Minimal hardware requirements & efficient translation

Segmentation can better support **sparse address spaces**

Avoids **internal fragmentation**.

Minimises memory waste between logical segments of the address space

Limitations of Segmentation

1) No expandable memory

Static memory allocation

2) No memory Sharing

Cannot share memory between processes

3) Non-Relative Memory Addresses

Location of code & data determined at runtime

4) External Fragmentation

Cannot relocate/move programs. Leads to fragmentation

5) Internal Fragmentation

Address Space must be contiguous

Segmentation Summary Cons

External fragmentation still a problem
Must fit variable-sized chunks into physical memory.

May move processes multiple times to fit everything

